

Apache Kafka



Anderson Giuseppe Saraiva Patriarca Fantin - 9016981 Gustavo Câmara Reis - 11207780 Laís Garcia - 11270628 Larissa Fabião da Fonseca - 11208367

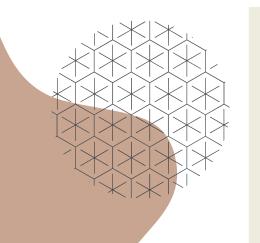
Agenda

Ol O que é o O3 Extensões Apache Kafka

O2 Como o Kafka O4 Exemplo funciona? Prático

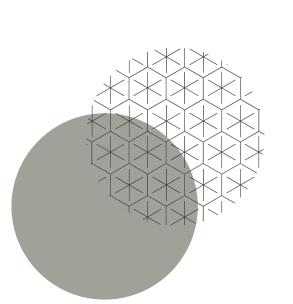


O que é o Apache Kafka



Plataforma open-source distribuída de streaming

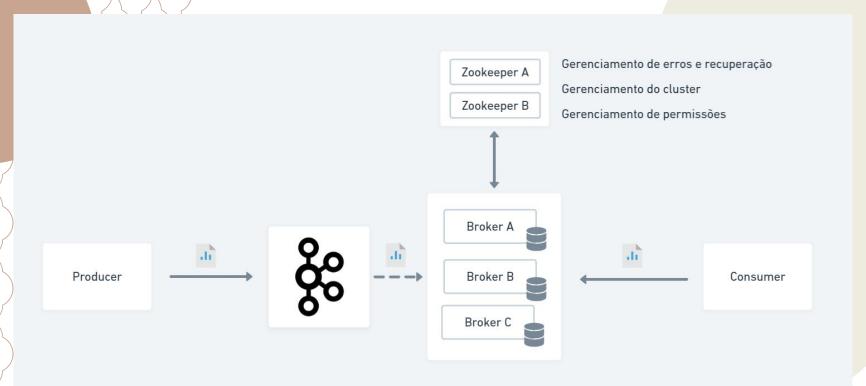
- Scala e Java
- Objetivo
- Fluxo básico com mensagens
- Qual a utilidade?
- Proporciona diversas maneiras de mover dados



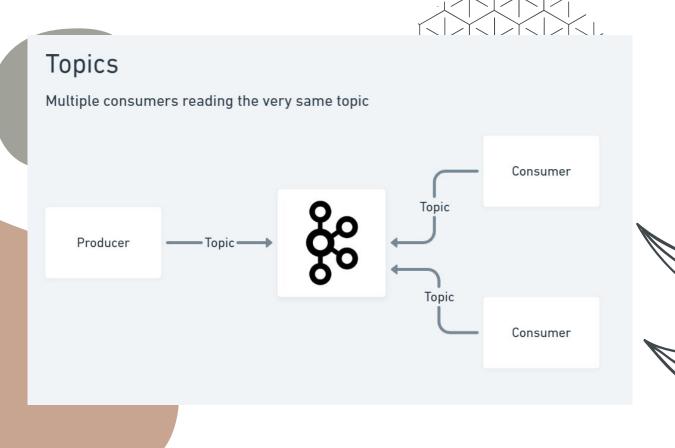
O2 Como funciona?



Ciclo geral

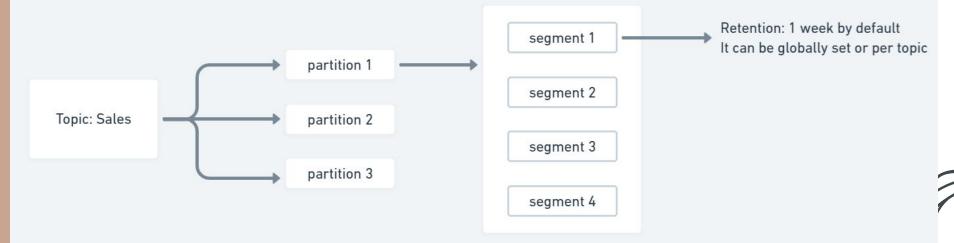


Armazenamento de informações



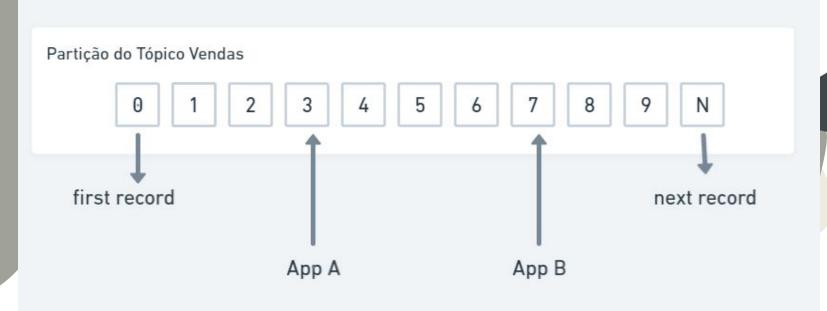
Armazenamento de informações





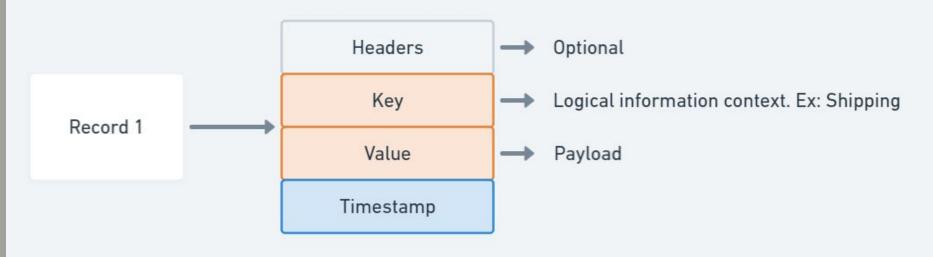
Leitura de informações

Topic log





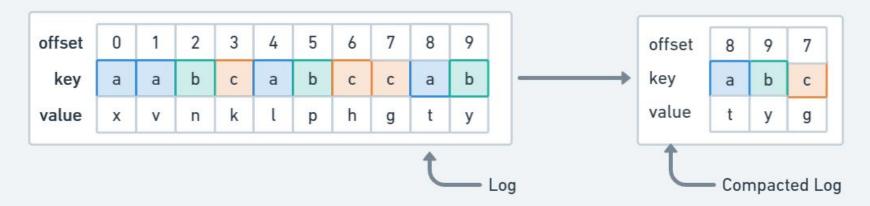
Record anatomy

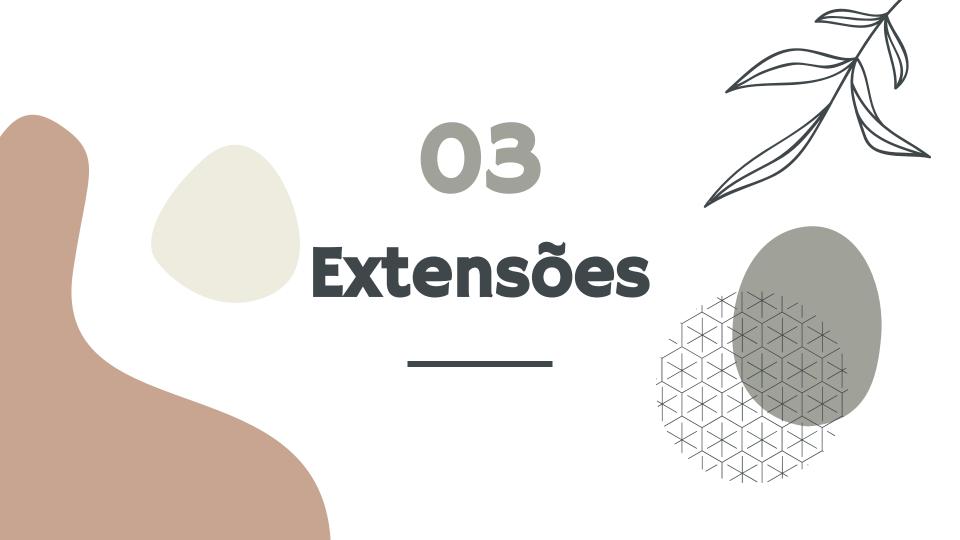


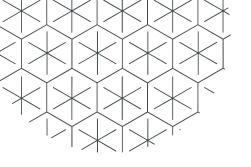


Compacted topic

Topic: Sales







Extensões



Connect

Processa e analiza dados armazenados no Kafka



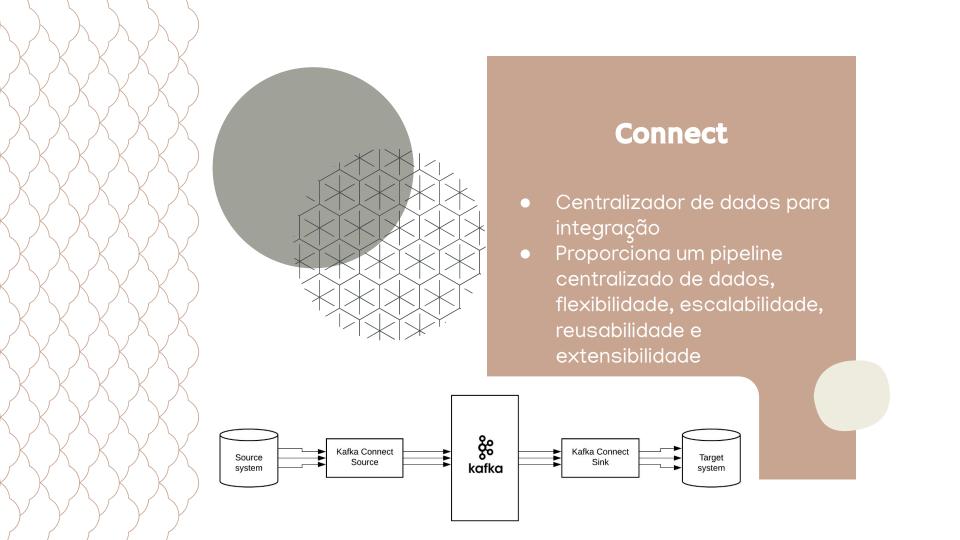
Streams

Framework para tratar, transformar e agregar as mensagens



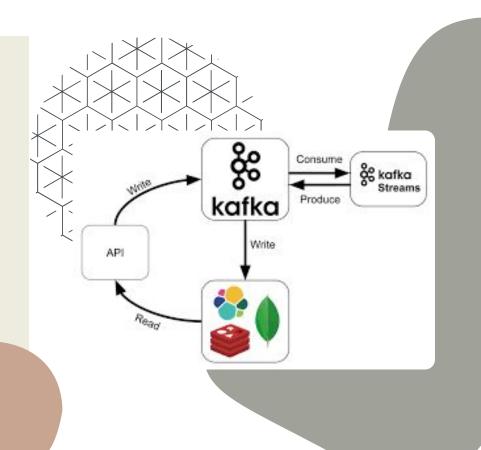
Schema Registry

Biblioteca de cliente onde os dados são armazenados em um cluster Kafka



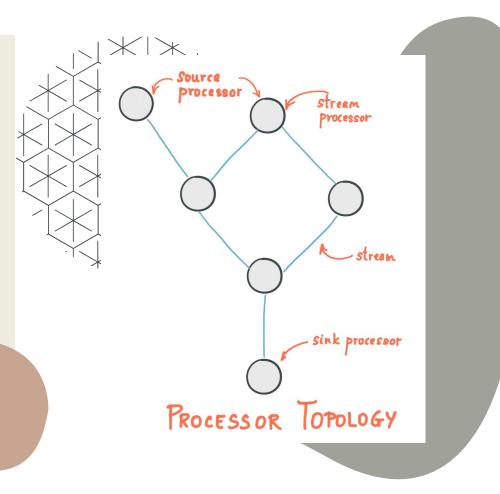
Streams

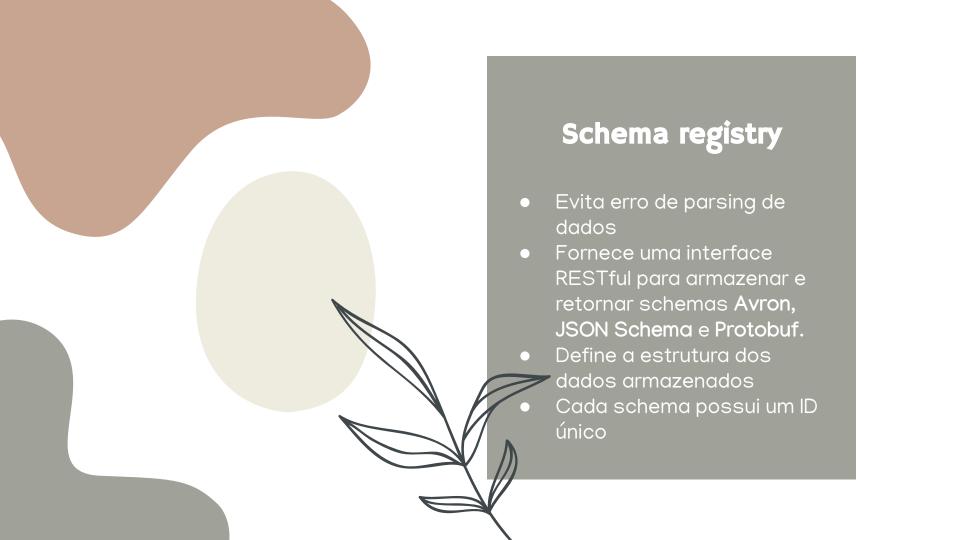
- API Funcional de Java
- Permite a criação de filtros, agregações, junções entre outras operações
- Escalável
- Tolerante a falhas
- Se integra como uma biblioteca

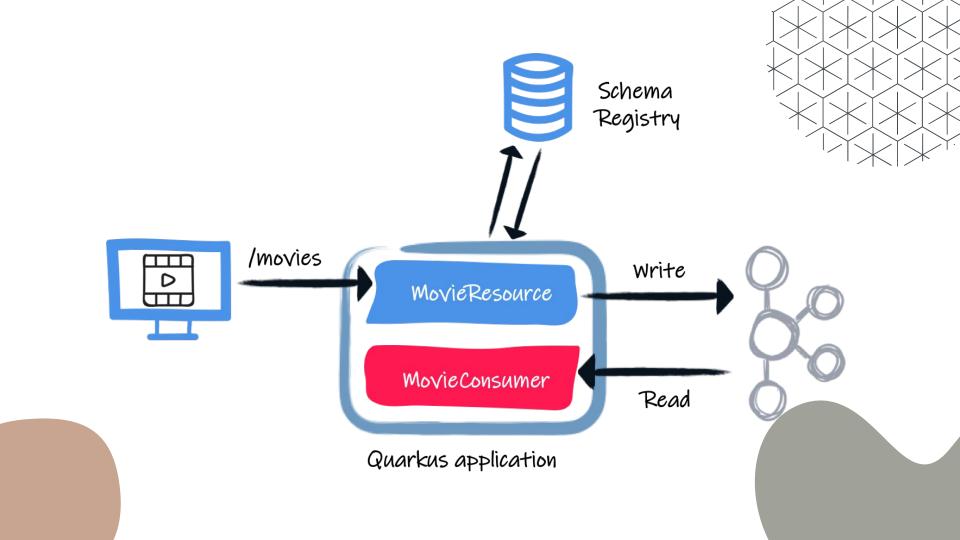


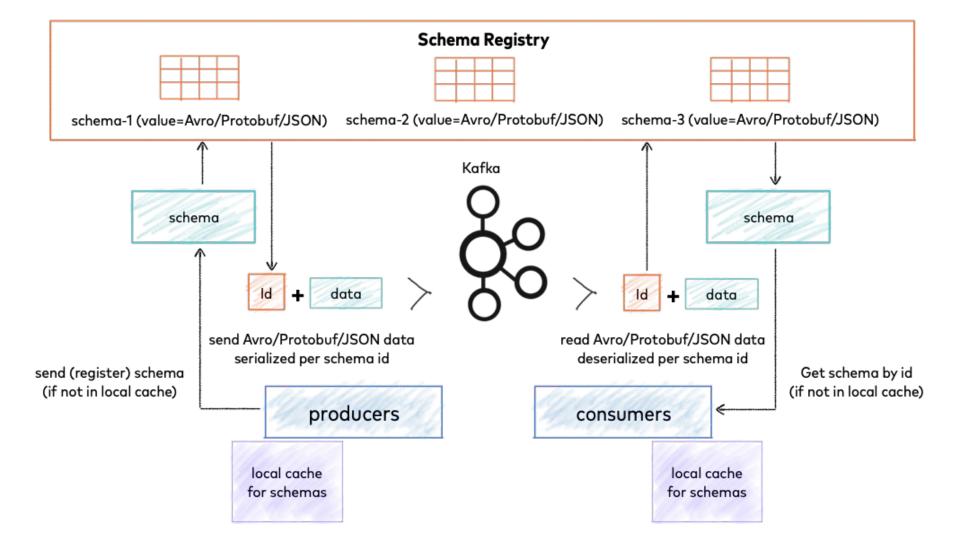
Streams

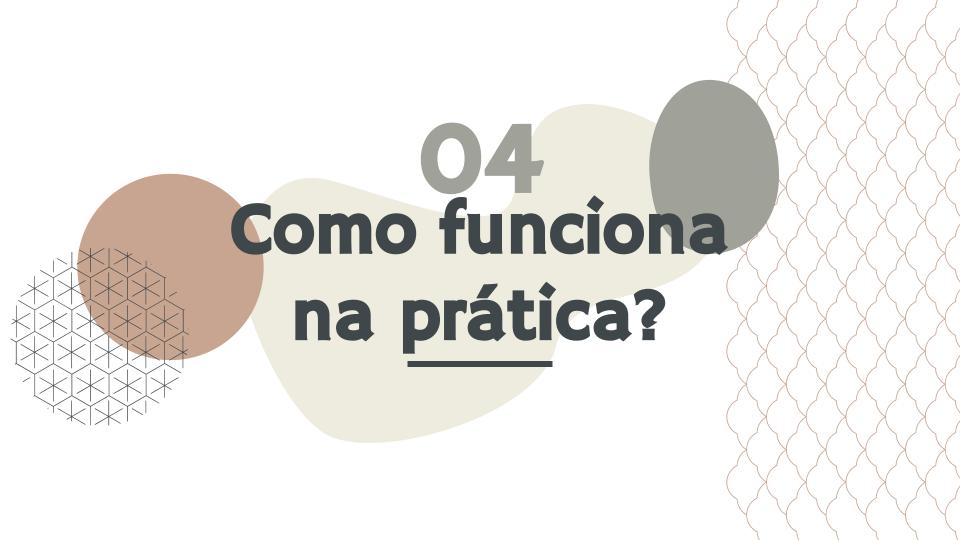
- Stream: uma sequência ordenada, reproduzível e tolerante a falhas de registros de dados
- Stream processor: um step de processamento
- Source processor: produz um fluxo de entrada consumindo registros do Kafka
- Sink processor: envia os registros recebidos de seus processadores up-stream para um tópico Kafka.





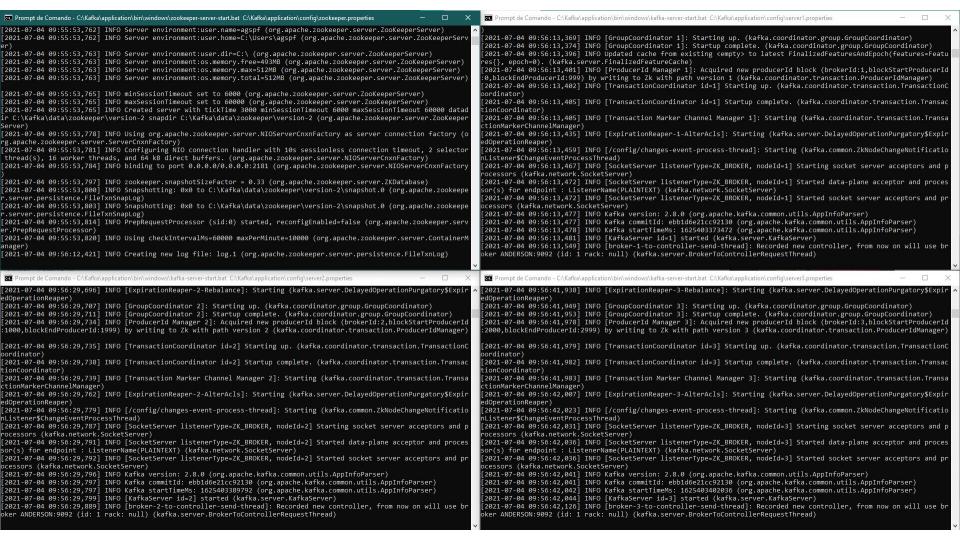






Inicialização dos serviços

- # Iniciar o zookeeper
- C:\Kafka\application\bin\windows\zookeeper-server-start.bat C:\Kafka\application\config\zookeeper.properties
- # Iniciar broker 1
- C:\Kafka\application\bin\windows\kafka-server-start.bat C:\Kafka\application\config\server1.properties
- # Iniciar broker 2
- C:\Kafka\application\bin\windows\kafka-server-start.bat C:\Kafka\application\config\server2.properties
- # Iniciar broker 3
- C:\Kafka\application\bin\windows\kafka-server-start.bat C:\Kafka\application\config\server3.properties



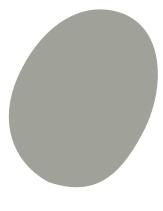


C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --create --topic topic-3-1 --partitions 3 --replication-factor 1 --bootstrap-server localhost:9092 Created topic topic-3-1.

C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-1 --bootstrap-server localhost:9092 Topic: topic-3-1 TopicId: NgZ1I_VRTLy9cyLGZlwJpw PartitionCount: 3 ReplicationFactor: 1 Configs: segment.bytes=1073741824

Topic: topic-3-1 Partition: 0 Leader: 3 Replicas: 3 Isr: 3
Topic: topic-3-1 Partition: 1 Leader: 1 Replicas: 1 Isr: 1
Topic: topic-3-1 Partition: 2 Leader: 2 Replicas: 2 Isr: 2







```
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --create --topic topic-3-1 --partitions 3 --replication-factor 1 --bootstrap-server localhost:9092
Created topic topic-3-1.
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-1 --bootstrap-server localhost:9092
Topic: topic-3-1
                       TopicId: NgZ1I VRTLy9cyLGZlwJpw PartitionCount: 3
                                                                              ReplicationFactor: 1 Configs: segment.bytes=1073741824
                              Partition: 0
       Topic: topic-3-1
                                              Leader: 3
                                                              Replicas: 3
                                                                              Isr: 3
                                              Leader: 1
       Topic: topic-3-1
                              Partition: 1
                                                              Replicas: 1
                                                                              Isr: 1
```

Isr: 2

Replicas: 1,3,2 Isr: 1,3,2

Replicas: 2,1,3 Isr: 2,1,3

Replicas: 3,2,1 Isr: 3,2,1

Replicas: 2

Topic: topic-3-1

Topic: topic-3-3

Topic: topic-3-3

Topic: topic-3-3

Partition: 2 Leader: 2

Partition: 2 Leader: 3

Leader: 1

Leader: 2

Partition: 0

Partition: 1

```
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092
                     TopicId: 8gyvzpwzRraVJW9Q-fB Ug PartitionCount: 3
Topic: topic-3-3
                                                                         ReplicationFactor: 3
                                                                                               Configs: segment.bytes=1073741824
       Topic: topic-3-3
                             Partition: 0
                                           Leader: 1
                                                          Replicas: 1,3,2 Isr: 1,3,2
                            Partition: 1 Leader: 2
      Topic: topic-3-3
Topic: topic-3-3
                                                          Replicas: 2,1,3 Isr: 2,1,3
                             Partition: 2 Leader: 3
                                                          Replicas: 3,2,1 Isr: 3,2,1
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092
Topic: topic-3-3
                        TopicId: 8gyvzpwzRraVJW9Q-fB_Ug PartitionCount: 3
                                                                                 ReplicationFactor: 3
                                                                                                         Configs: segment.bytes=1073741824
                                Partition: 0
                                                Leader: 1
                                                                Replicas: 1,3,2 Isr: 1,2
       Topic: topic-3-3
                                                                Replicas: 2,1,3 Isr: 2,1
       Topic: topic-3-3
                                Partition: 1
                                                Leader: 2
       Topic: topic-3-3
                                Partition: 2
                                                Leader: 2
                                                                 Replicas: 3,2,1 Isr: 2,1
```

```
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092
Topic: topic-3-3
                     TopicId: 8gyvzpwzRraVJW9Q-fB Ug PartitionCount: 3
                                                                        ReplicationFactor: 3
                                                                                             Configs: segment.bytes=1073741824
      Topic: topic-3-3
                            Partition: 0
                                          Leader: 1
                                                         Replicas: 1,3,2 Isr: 1,3,2
      Topic: topic-3-3
                            Partition: 1
                                          Leader: 2
                                                         Replicas: 2,1,3 Isr: 2,1,3
      Topic: topic-3-3
                                                         Replicas: 3,2,1 Isr: 3,2,1
                            Partition: 2 Leader: 3
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092
Topic: topic-3-3
                       TopicId: 8gyvzpwzRraVJW9Q-fB_Ug PartitionCount: 3
                                                                                ReplicationFactor: 3
                                                                                                        Configs: segment.bytes=1073741824
                               Partition: 0
                                               Leader: 1
                                                                Replicas: 1,3,2 Isr: 1,2
       Topic: topic-3-3
                                                               Replicas: 2,1,3 Isr: 2,1
       Topic: topic-3-3
                               Partition: 1
                                               Leader: 2
       Topic: topic-3-3
                               Partition: 2
                                               Leader: 2
                                                                Replicas: 3,2,1 Isr: 2,1
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092
Topic: topic-3-3
                       TopicId: 8gyvzpwzRraVJW9Q-fB Ug PartitionCount: 3
                                                                               ReplicationFactor: 3
                                                                                                       Configs: segment.bytes=1073741824
       Topic: topic-3-3
                               Partition: 0
                                               Leader: 1
                                                               Replicas: 1,3,2 Isr: 1,2,3
       Topic: topic-3-3
                               Partition: 1
                                                               Replicas: 2,1,3 Isr: 2,1,3
                                               Leader: 2
       Topic: topic-3-3
                               Partition: 2
                                                               Replicas: 3,2,1 Isr: 2,1,3
                                               Leader: 2
```

C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092

Leader: 3

Partition: 2

Topic: topic-3-3

```
Configs: segment.bytes=1073741824
Topic: topic-3-3
                     TopicId: 8gyvzpwzRraVJW9Q-fB Ug PartitionCount: 3
                                                                       ReplicationFactor: 3
      Topic: topic-3-3
                            Partition: 0
                                          Leader: 1
                                                         Replicas: 1,3,2 Isr: 1,3,2
                            Partition: 1
                                                         Replicas: 2,1,3 Isr: 2,1,3
      Topic: topic-3-3
                                          Leader: 2
      Topic: topic-3-3
                                                         Replicas: 3,2,1 Isr: 3,2,1
                            Partition: 2
                                          Leader: 3
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092
Topic: topic-3-3
                       TopicId: 8gyvzpwzRraVJW9Q-fB Ug PartitionCount: 3
                                                                               ReplicationFactor: 3
                                                                                                       Configs: segment.bytes=1073741824
                               Partition: 0
                                               Leader: 1
                                                               Replicas: 1,3,2 Isr: 1,2
       Topic: topic-3-3
                                                               Replicas: 2,1,3 Isr: 2,1
       Topic: topic-3-3
                               Partition: 1
                                               Leader: 2
       Topic: topic-3-3
                               Partition: 2
                                               Leader: 2
                                                               Replicas: 3,2,1 Isr: 2,1
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092
Topic: topic-3-3
                       TopicId: 8gyvzpwzRraVJW9Q-fB Ug PartitionCount: 3
                                                                               ReplicationFactor: 3
                                                                                                       Configs: segment.bytes=1073741824
       Topic: topic-3-3
                               Partition: 0
                                               Leader: 1
                                                               Replicas: 1,3,2 Isr: 1,2,3
       Topic: topic-3-3
                               Partition: 1
                                                               Replicas: 2,1,3 Isr: 2,1,3
                                               Leader: 2
       Topic: topic-3-3
                               Partition: 2
                                                               Replicas: 3,2,1 Isr: 2,1,3
                                               Leader: 2
C:\>C:\Kafka\application\bin\windows\kafka-topics.bat --describe --topic topic-3-3 --bootstrap-server localhost:9092
                       TopicId: 8gyvzpwzRraVJW9Q-fB Ug PartitionCount: 3
                                                                                                       Configs: segment.bytes=1073741824
Topic: topic-3-3
                                                                               ReplicationFactor: 3
       Topic: topic-3-3
                               Partition: 0
                                               Leader: 1
                                                               Replicas: 1,3,2 Isr: 1,2,3
       Topic: topic-3-3
                               Partition: 1
                                               Leader: 2
                                                               Replicas: 2,1,3 Isr: 2,1,3
```

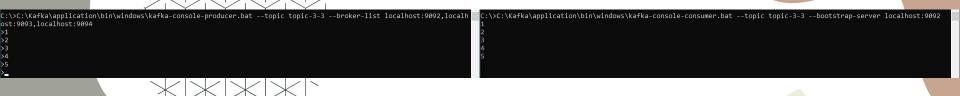
Replicas: 3,2,1 Isr: 2,1,3



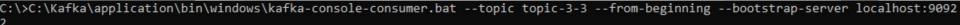
C:\>C:\>C:\Kafka\application\bin\windows\kafka-console-producer.bat --topic topic-3-3 --broker-list localhost:9092,localh
ost:9093,localhost:9094
>1
2
>2
>3
>3
>4
>4
>5



C:\>C:\Kafka\application\bin\windows\kafka-console-consumer.bat --topic topic-3-3 --bootstrap-server localhost:9092



C:\>C:\Kafka\application\bin\windows\kafka-console-consumer.bat --topic topic-3-3 --bootstrap-server localhost:9092





Produtor e consumidores no mesmo grupo



C:\>C:\Kafka\application\bin\windows\kafka-consumer-groups.bat --group group-1 --describe --bootstrap-server localhost:9092

| GROUP group-1 group-1 group-1 | TOPIC topic-3-3 topic-3-3 topic-3-3 | PARTITION 0 1 2 | CURRENT-OFFSET 9 5 7 | LOG-END-OFFSET 9 5 7 | LAG 0 0 | CONSUMER-ID HOST consumer-group-1-1-6ec31caf-d44e-4325-9452-26c07328492c /192.168.0.1/consumer-group-1-1-6ec31caf-d44e-4325-9452-26c07328492c /192.168.0.1/consumer-group-1-1-d3f3d6ae-b82c-4248-9238-18b222a9688a /192.168.0.1/c | consumer-group- |
|--|--|--------------------------|-------------------------------|-------------------------------|---------------|---|-----------------|
| group-1 | topic-3-3 | 2 | 7 | 7 | 0 | consumer-group-1-1-d3f3d6ae-b82c-4248-9238-18b222a9688a /192.168.0.1 | consumer-group- |

Produtor e consumidores no mesmo grupo

C:\>C:\Kafka\application\bin\windows\kafka-consumer-groups.bat --group group-1 --describe --bootstrap-server localhost:9092

GROUP TOPIC PARTITION CURRENT-OFFSET LOG-END-OFFSET LAG CONSUMER-ID HOST CLIENT-ID group-1 topic-3-3 0 9 0 consumer-group-1-1-6ec31caf-d44e-4325-9452-26c07328492c /192.168.0.14 consumer-group-1-1 group-1 topic-3-3 1 5 5 0 consumer-group-1-1-6ec31caf-d44e-4325-9452-26c07328492c /192.168.0.14 consumer-group-1-1

consumer-group-1-1-d3f3d6ae-b82c-4248-9238-18b222a9688a /192.168.0.14

consumer-group-1-1



topic-3-3

group-1

| group-1 t | opic-3-3 | 1 | 5 | LOG-END-OFFSET 5 | 0 | CONSUMER-ID HOST consumer-group-1-1-c3153fd8-ea31-49f1-a546-a2fad4f318b6 /192. | 0212111 20 |
|-----------|----------|---|---|---------------------|---|--|----------------------------|
| | opic-3-3 | 2 | 7 | 7 | 0 | consumer-group-1-1-d3f3d6ae-b82c-4248-9238-18b222a9688a /192. | |
| group-1 t | opic-3-3 | 0 | 9 | 9 | 0 | consumer-group-1-1-6ec31caf-d44e-4325-9452-26c07328492c /192. | .168.0.14 consumer-group-1 |

```
Produtor em Java
```

Foublic class ProducerCreator {

31

```
public static String KAFKA BROKERS = "localhost:9092";
          public static String TOPIC NAME = "topic-3-3";
          public static int MESSAGE COUNT = 1000;
          public static Producer<Long, String> createProducer() {
              Properties props = new Properties();
              props.put(ProducerConfig.BOOTSTRAP SERVERS CONFIG, KAFKA BROKERS);
              props.put(ProducerConfig.KEY SERIALIZER CLASS CONFIG, LongSerializer.class.getName());
              props.put(ProducerConfig.VALUE SERIALIZER CLASS CONFIG, StringSerializer.class.getName());
              return new KafkaProducer<>(props);
12
13
14
    -static void runProducer() {
          Producer<Long, String> producer = ProducerCreator.createProducer():
16
17
          for (int index = 0; index < IKafkaConstants.MESSAGE COUNT; index++) {
18
              ProducerRecord<Long, String> record = new ProducerRecord<Long, String>(ProducerCreator.TOPIC NAME, "This is record " + index);
19
              try {
                  RecordMetadata metadata = producer.send(record).get();
                  System.out.println("Record sent with key " + index + " to partition " + metadata.partition() + " with offset " + metadata.offset());
              } catch (ExecutionException e) {
23
                  System.out.println("Error in sending record");
                  System.out.println(e);
25
26
               catch (InterruptedException e) {
                  System.out.println("Error in sending record");
27
                  System.out.println(e);
```

Consumidor em Java

```
Public class ConsumerCreator {
          public static String KAFKA BROKERS = "localhost:9092";
          public static String GROUP ID = "group-1";
          public static String TOPIC NAME = "topic-3-3";
          public static Consumer < Long, String> createConsumer() {
              Properties props = new Properties():
              props.put(ConsumerConfig.BOOTSTRAP SERVERS CONFIG, KAFKA BROKERS);
              props.put(ConsumerConfig.GROUP ID CONFIG, GROUP ID);
              props.put(ConsumerConfig.KEY DESERIALIZER CLASS CONFIG, LongDeserializer.class.getName());
              props.put(ConsumerConfig.VALUE DESERIALIZER CLASS CONFIG, StringDeserializer.class.getName());
12
              Consumer < Long, String > consumer = new KafkaConsumer <> (props);
13
              consumer.subscribe(Collections.singletonList(TOPIC NAME));
14
              return consumer;
15
16
    Fistatic void runConsumer() {
19
          Consumer<Long, String> consumer = ConsumerCreator.createConsumer();
20
          int noMessageFound = 0;
21
          while (true) {
22
            ConsumerRecords<Long, String> consumerRecords = consumer.poll(1000);
23
            if (consumerRecords.count() == 0) {
24
                noMessageFound++:
25
                if (noMessageFound > IKafkaConstants.MAX NO MESSAGE FOUND COUNT)
26
                  break:
27
            consumerRecords.forEach(record -> {
28
29
                System.out.println("Record Key " + record.key());
30
                System.out.println("Record value " + record.value());
                System.out.println("Record partition " + record.partition());
31
32
                System.out.println("Record offset " + record.offset());
33
34
             consumer.commitAsync();
35
36
          consumer.close();
37
```



Referências

- https://docs.confluent.io/platform/current/schema-registry/index.html
- https://kafka.apache.org/documentation/
- https://kafka.apache.org/quickstart